

ECE 350
Real-time
Operating
Systems



Lecture 12: Reliable Storage

Prof. Seyed Majid Zahedi

<https://ece.uwaterloo.ca/~smzahedi>

Outline

- Problem posed by machine/disk failures
- Transaction concept (ASID properties)
- Transactional file systems
 - Journaling and logging
- Copy-on-write file systems
- Redundant arrays of independent disks (RAID)

Recall: Important “abilities”

- **Reliability**: ability of system or component to perform its required functions under stated conditions for specified time
 - System is not only “up”, but also working correctly
 - Includes availability, security, fault tolerance/durability
 - Data must survive system crashes, disk crashes, other problems
- **Availability**: probability that system can accept and process requests
 - Can build highly-available systems, despite unreliable components
 - Involves independence of failures and redundancy
 - Often as “nines” of probability. 99.9% is “3-nines of availability”
- **Durability**: ability of system to recover data despite faults
 - This idea is fault-tolerance applied to data
 - Doesn't necessarily imply availability: information on pyramids was very durable, but could not be accessed until discovery of Rosetta Stone

File System Reliability

- Single logical file operation can involve updates to **multiple** physical storage blocks
 - inode, indirect block, data block, bitmap, ...
 - With sector/page remapping, single update to physical storage block can require multiple (even lower-level) updates to sectors/pages
- What can happen if disk loses power or software **crashes**?
 - Some operations in progress may complete
 - Some operations in progress may be lost
 - Overwrite of block may only partially complete
- How do we guarantee consistency regardless of when crash occurs?

Reliability Take I: Careful Ordering

- Sequence operations in specific order
 - Careful design to allow sequence to be interrupted safely
- Post-crash recovery
 - Read data structures to see if there were any operations in progress
 - Clean up/finish as needed
- Approach taken by
 - FAT and FFS (**fsck**) to protect file-system structure/metadata
 - Many app-level recovery schemes (e.g., Word, emacs autosaves)

Append Data to File in FAT

Normal operation

- Allocate data block
- Write data
- Write new MFT entry to point to data block
- Update file tail to point to new MFT entry
- Update access time at head of file

Recovery

- Scan MFT
- If entry is unlinked, delete data block
- If access time is incorrect, update

Create New File in FFS

Normal operation

- Allocate data block
- Write data block
- Allocate inode
- Write inode block
- Update bitmap of free blocks and inodes
- Update directory with file name
→ inode number
- Update modify time for directory

Recovery

- Scan inode table
- If any unlinked files (not in any directory), delete or put in lost & found directory
- Compare free block bitmap against inode trees
- Scan directories for missing update/access times
- Recovery time is proportional to storage size

Some General Rules for Careful Ordering

- Never write a pointer before initializing the block it points to (e.g., indirect block)
- Never reuse resource (e.g., inode, disk block, etc.) before nullifying all pointers to it
- Never clear last pointer to resource before setting new one (e.g., `mv`)



Careful Ordering: Discussion

- + Works with minimal support from storage drive
- + Works for most multi-step operations
- – Time-consuming recovery after each failure
- – Slow updates due to sync barriers between dependent operations
- – Hard to turn every operation to safely-interruptible sequence of writes
- – Hard to achieve consistency when multiple operations run concurrently

Reliability Take 2: Transaction Concept

- Transaction is group of operations providing ACID properties
- Atomicity: operations appear to **logically** happen as group, or not at all
 - At physical level, only single disk/flash write is atomic
- Consistency: transactions maintain data **integrity**
 - Each transaction moves system from one legal state to another
 - E.g., balance cannot be negative
 - E.g., cannot reschedule meeting on February 30
- Isolation: each transaction is not **affected** by other in-progress transactions
 - If multiple transactions execute concurrently, for each pair of transactions A and B, it either appears that A executed entirely before B or vice versa
- Durability: if transaction commits, its effects **persist** despite crashes

Typical Structure

- **Begin** transaction – get transaction id
- Do bunch of updates
 - If any fail along the way, **roll-back**
 - Or, if any conflicts with other transactions, roll-back
- **Commit** the transaction

“Classic” Transaction Example

- Transfer \$100 from Tom's account to Mike's account

BEGIN;

UPDATE accounts SET balance = balance - 100.00 WHERE name = 'Tom';

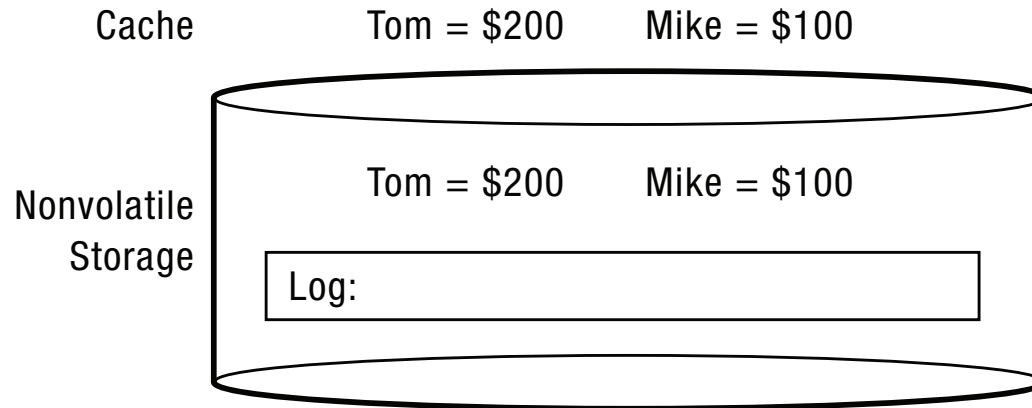
UPDATE accounts SET balance = balance + 100.00 WHERE name = 'Mike';

COMMIT;

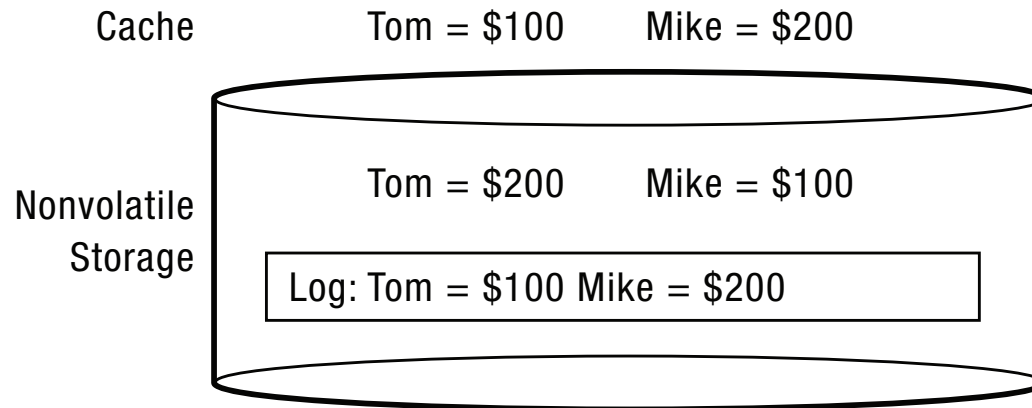
Atomicity and Durability: Redo Logging

- **Prepare**: append all updates to log
- **Commit**: append commit record to log
 - Single disk write to make transaction durable
- **Redo**: copy all updates to disk
- **Garbage collect**: reclaim space in log
- **Recover**: execute recovery routine if system crashes
 - Scan sequentially through log, do the following for each type of record
 - Update record: add to list of updates planned for this transaction
 - Commit record: redo all planned updates for this transaction
 - Roll-back record: discard list of updates planned for this transaction
 - At the end of log, discard updates for transactions without commit record
 - Garbage collect log

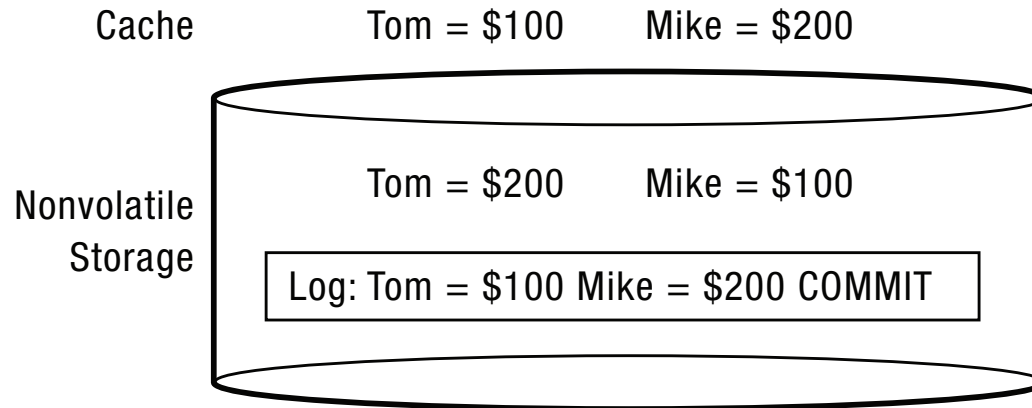
Redo Logging Example: Before Transaction Start



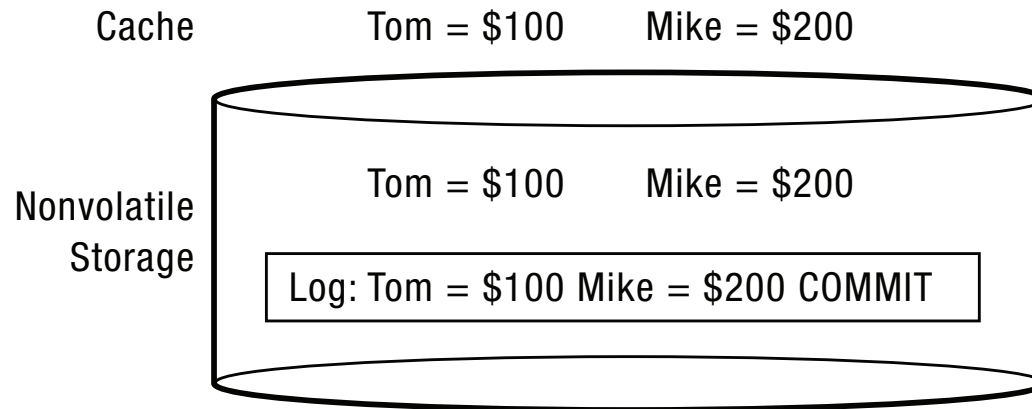
Redo Logging Example: After Updates Are Logged



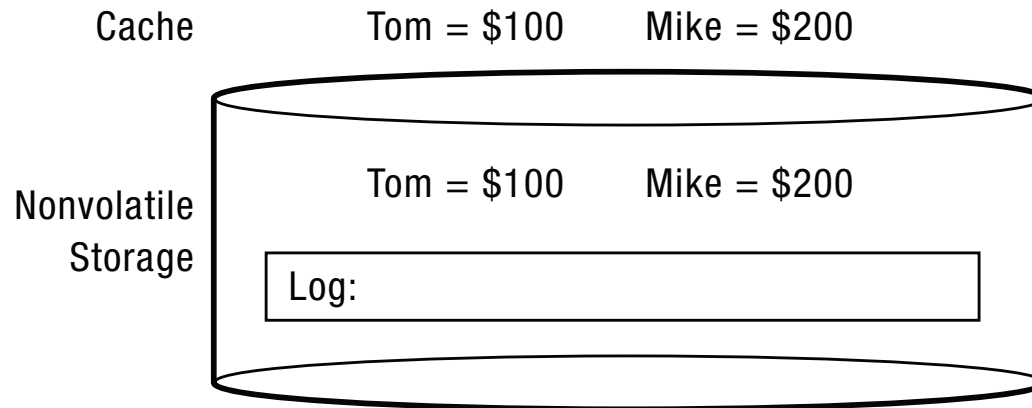
Redo Logging Example: After Commit Logged



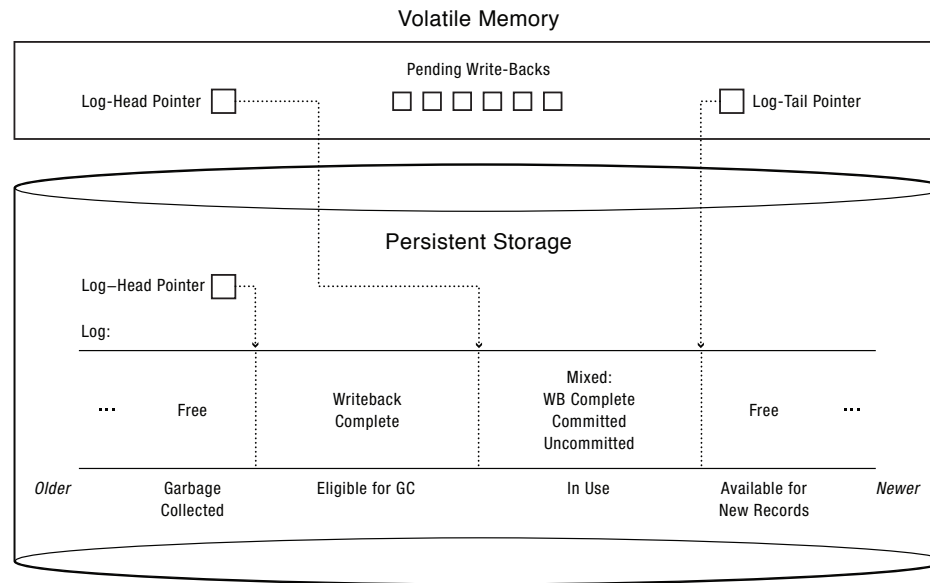
Redo Logging Example: After Redo



Redo Logging Example: After Garbage Collection



Redo Log Implementation



- Volatile memory has pointers to head and tail log
- New transaction records are appended to log's tail and cached in volatile memory
- Redo asynchronously writes pending updates for committed transactions
- Garbage collector periodically advances persistent log-head pointer so that recovery can skip at least some of committed transactions

Isolation: Two-phase Locking


- **Expanding phase**: locks may be acquired but not released
- **Contracting phase**: locks may be released but not acquired
- Transactions release locks only **AFTER** they commit
 - Prevents transactions from seeing results of non-committed ones
 - Could lead to **deadlock**
- If set of transactions deadlocks, one or more can be forced to **roll back**, release their locks, and **restart** at some **later** time
- **Serializability**: with two phase locking and redo logging, transactions appear to occur in sequential order

Transactional File Systems

- **Journaling**: apply updates to system metadata using transactions, but updates to users' files are done in place
 - E.g., NTFS, Apple HFS+, Linux XFS, JFS,
 - E.g., Ext3 and Ext4 use journaling by default
- **Logging**: apply all updates (metadata & data) using transactions
 - E.g., Ext3 and Ext4 can be configured to use logging

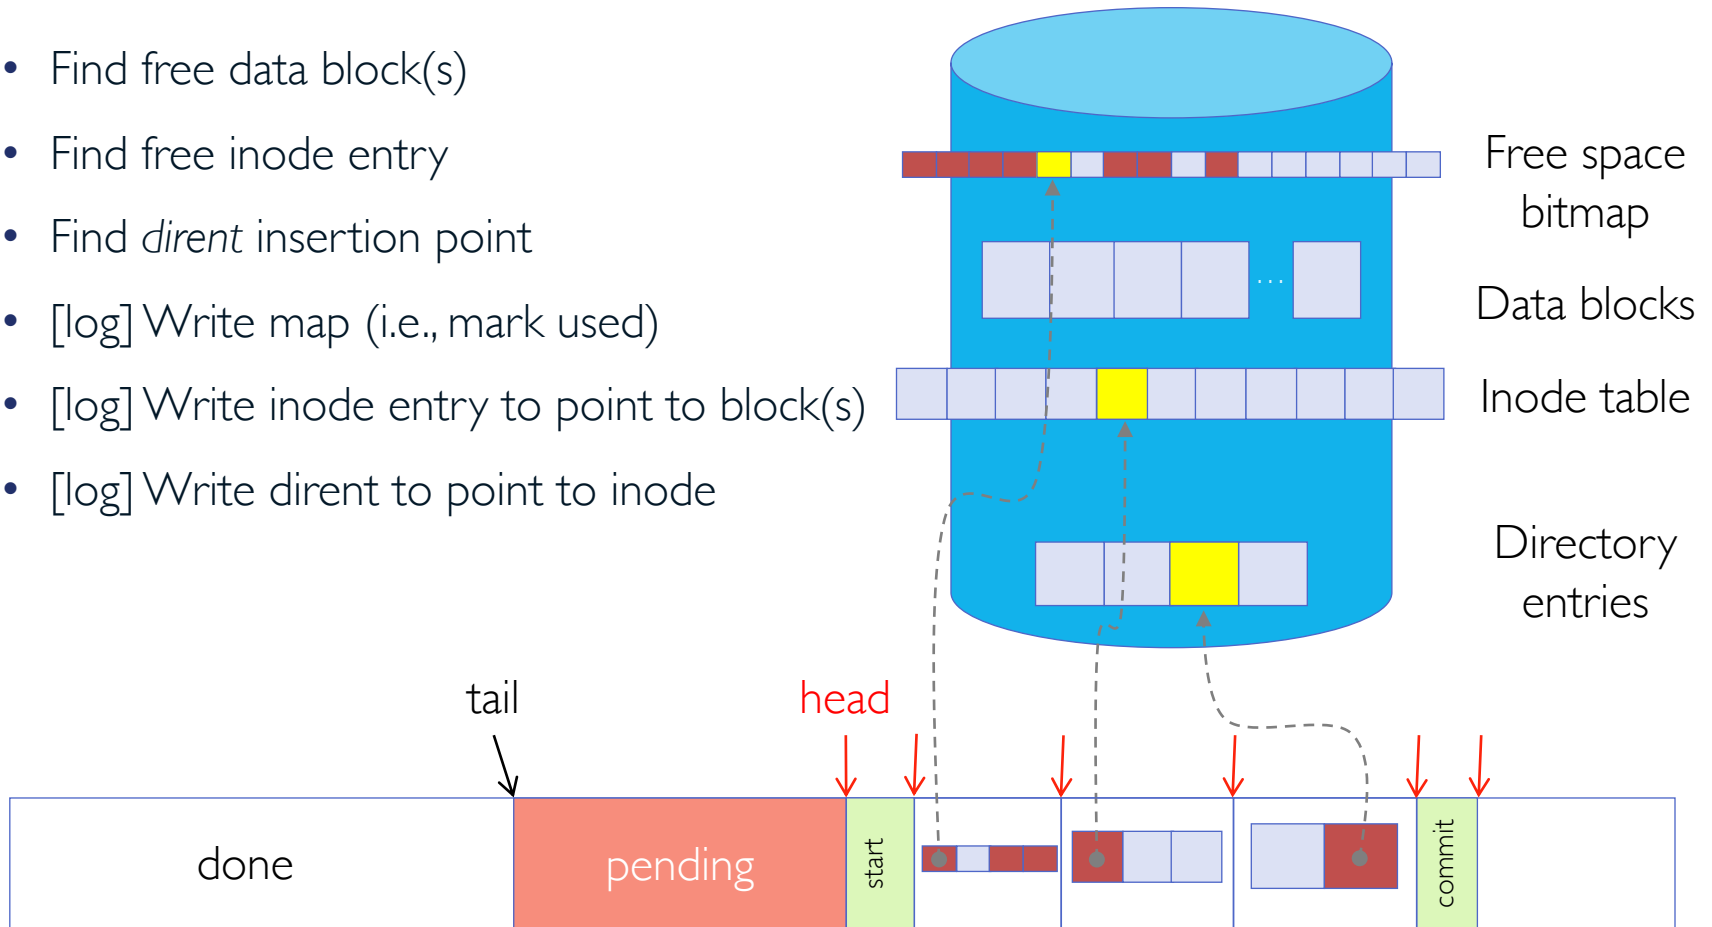
Example: Creating File

- Find free data block(s)
- Find free inode entry
- Find *dirent* insertion point
- Write map (i.e., mark used)
- Write inode entry to point to block(s)
- Write dirent to point to inode



Example: Creating File (as Transaction)


- Find free data block(s)
- Find free inode entry
- Find *dirent* insertion point
- [log] Write map (i.e., mark used)
- [log] Write inode entry to point to block(s)
- [log] Write dirent to point to inode



Log in non-volatile storage (Flash or on Disk)

Example: Creating File (as Transaction)

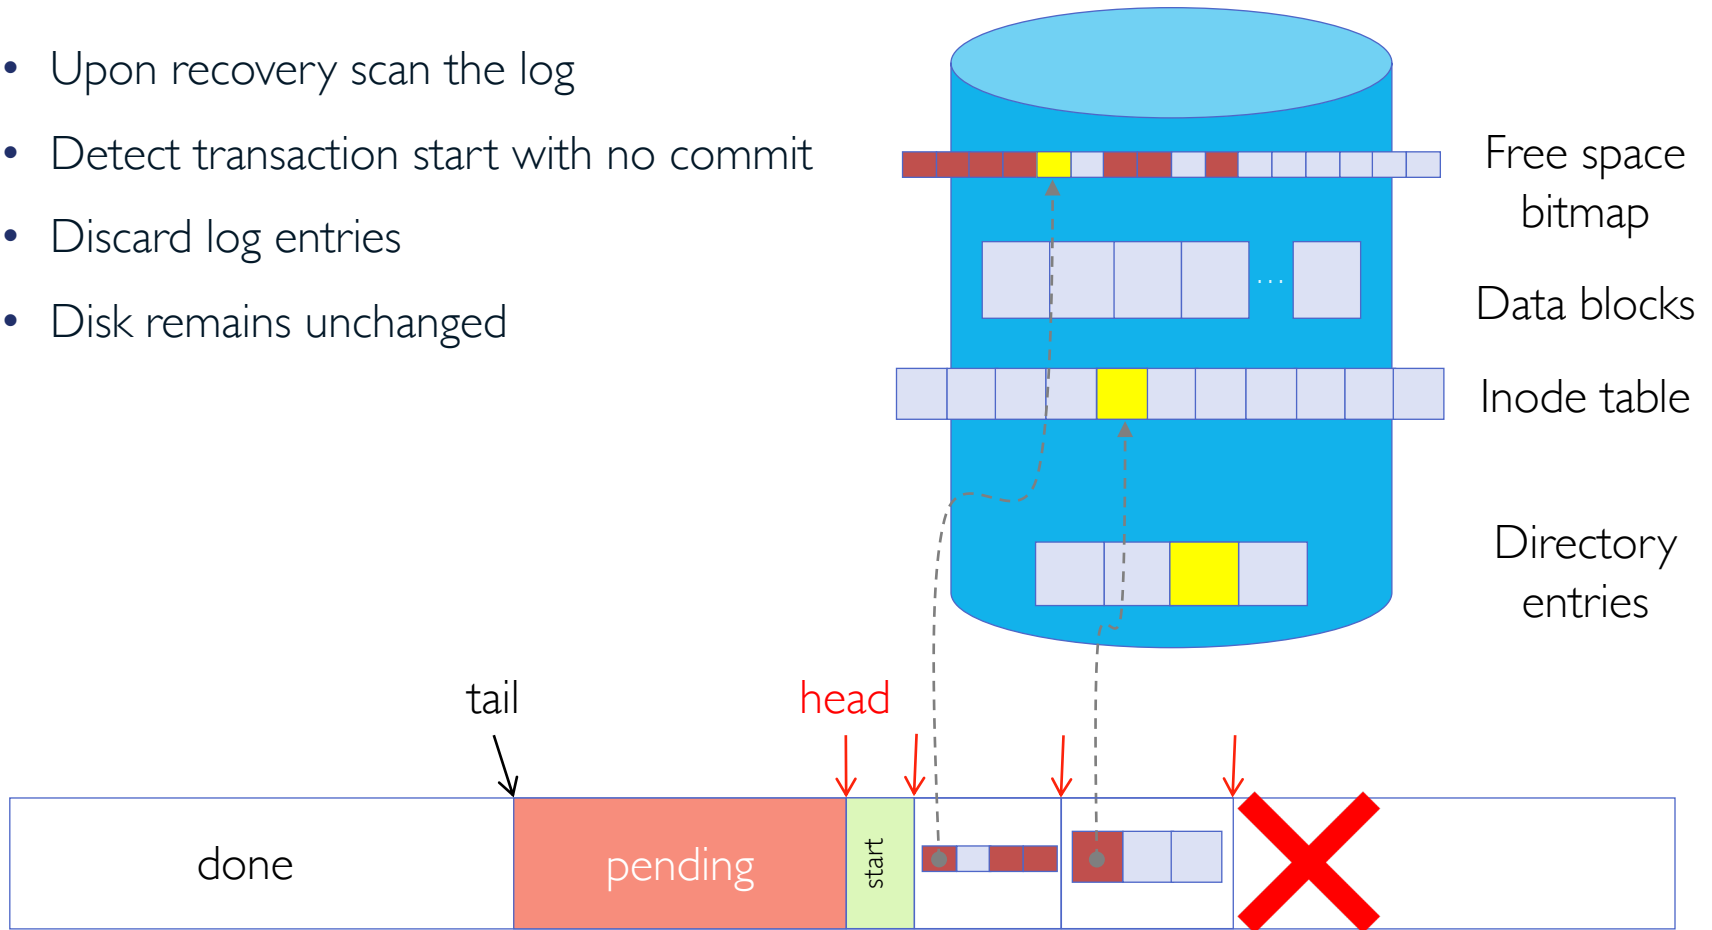
- Find free data block(s)
- Find free inode entry
- Find *dirent* insertion point
- [log] Write map (i.e., mark used)
- [log] Write inode entry to point to block(s)
- [log] Write dirent to point to inode



Log in non-volatile storage (Flash or on Disk)

Example: Recovery After Failure During Logging

- Upon recovery scan the log
- Detect transaction start with no commit
- Discard log entries
- Disk remains unchanged

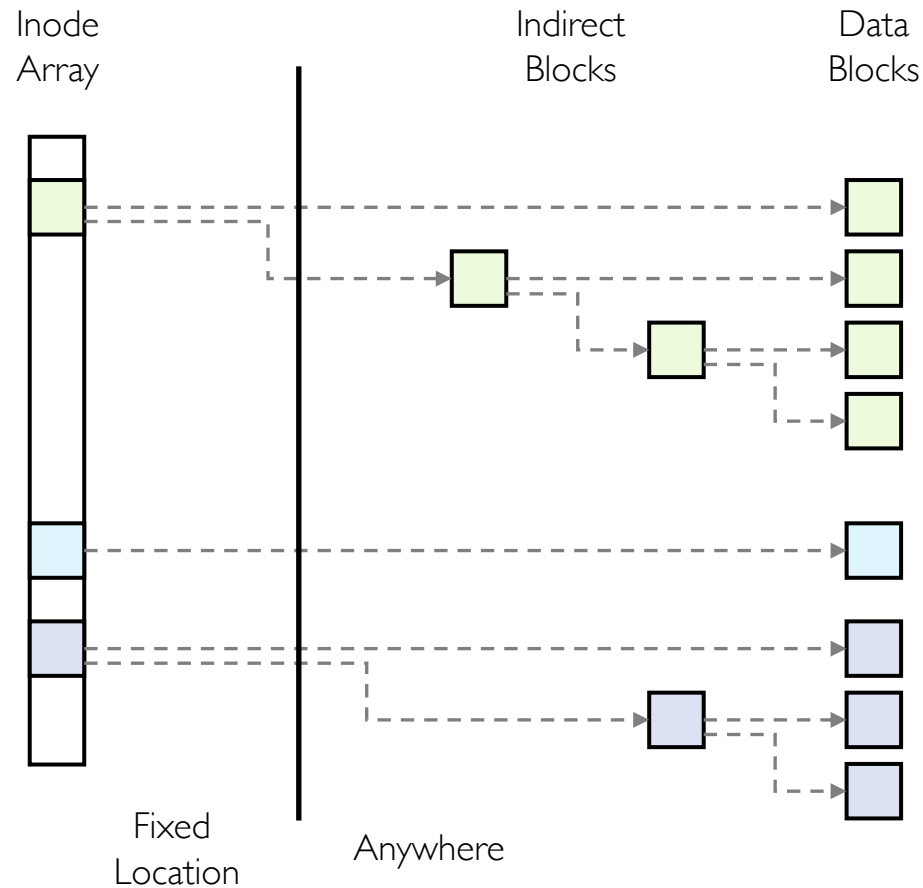


Log in non-volatile storage (Flash or on Disk)

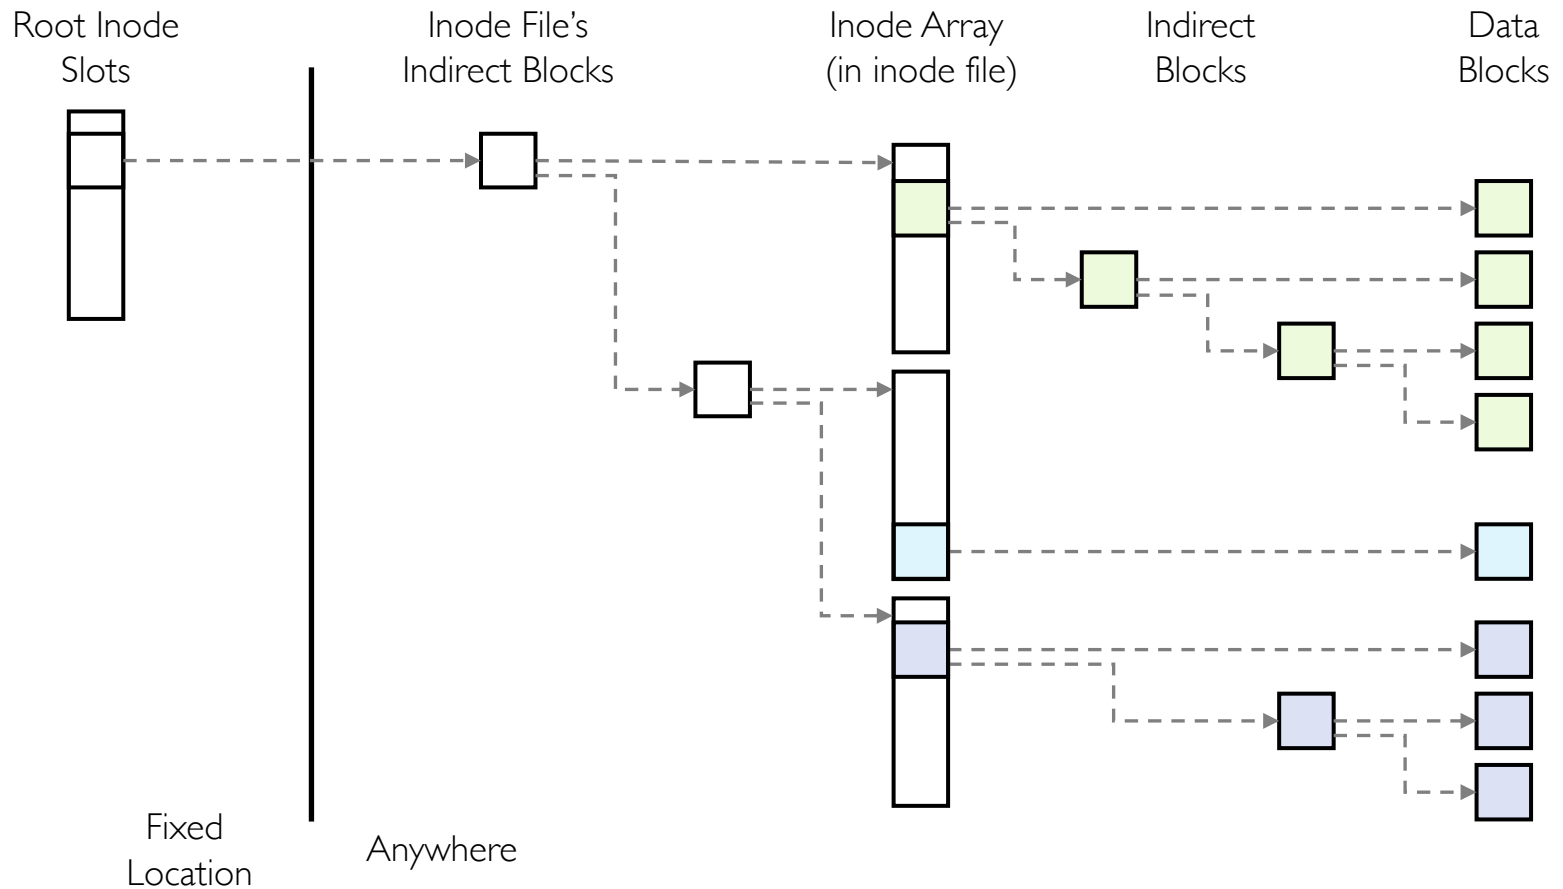
Copy-on-write (COW) File System

- For each update, write new version of file system
 - Never update in place
 - Reuse existing unchanged disk blocks
- Seems expensive! But ...
 - Updates can be batched
 - Almost all disk writes can occur in parallel
- Approach taken in network file server appliances
 - NetApp's Write Anywhere File Layout (WAFL)
 - ZFS (Sun/Oracle) and OpenZFS

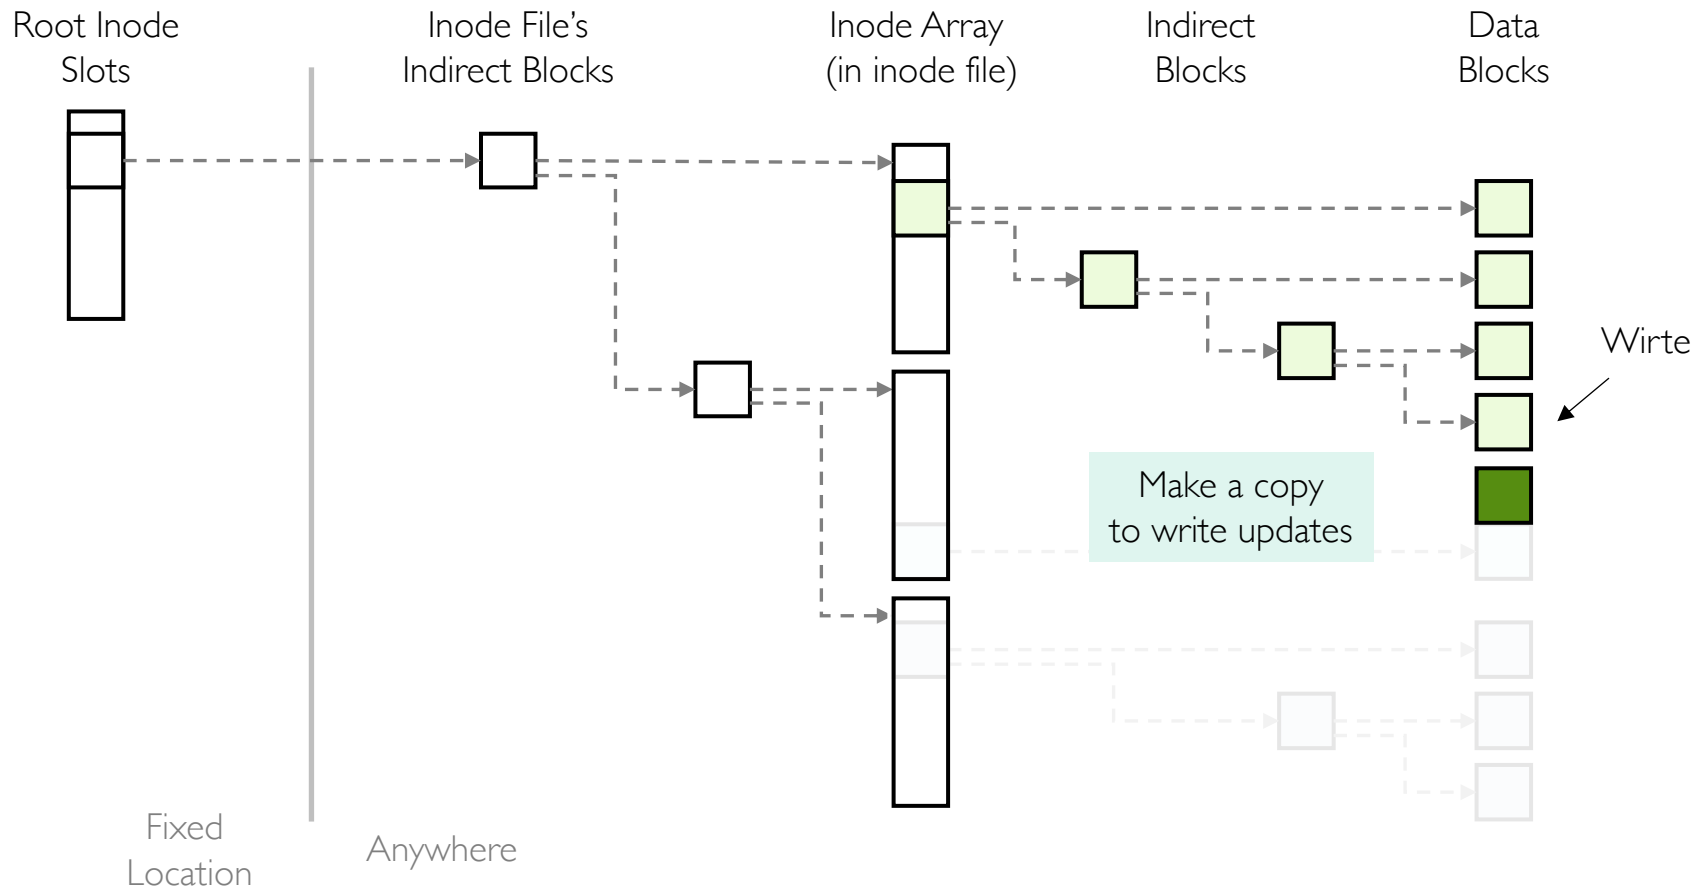
Recall: Traditional, Update-in-place File System



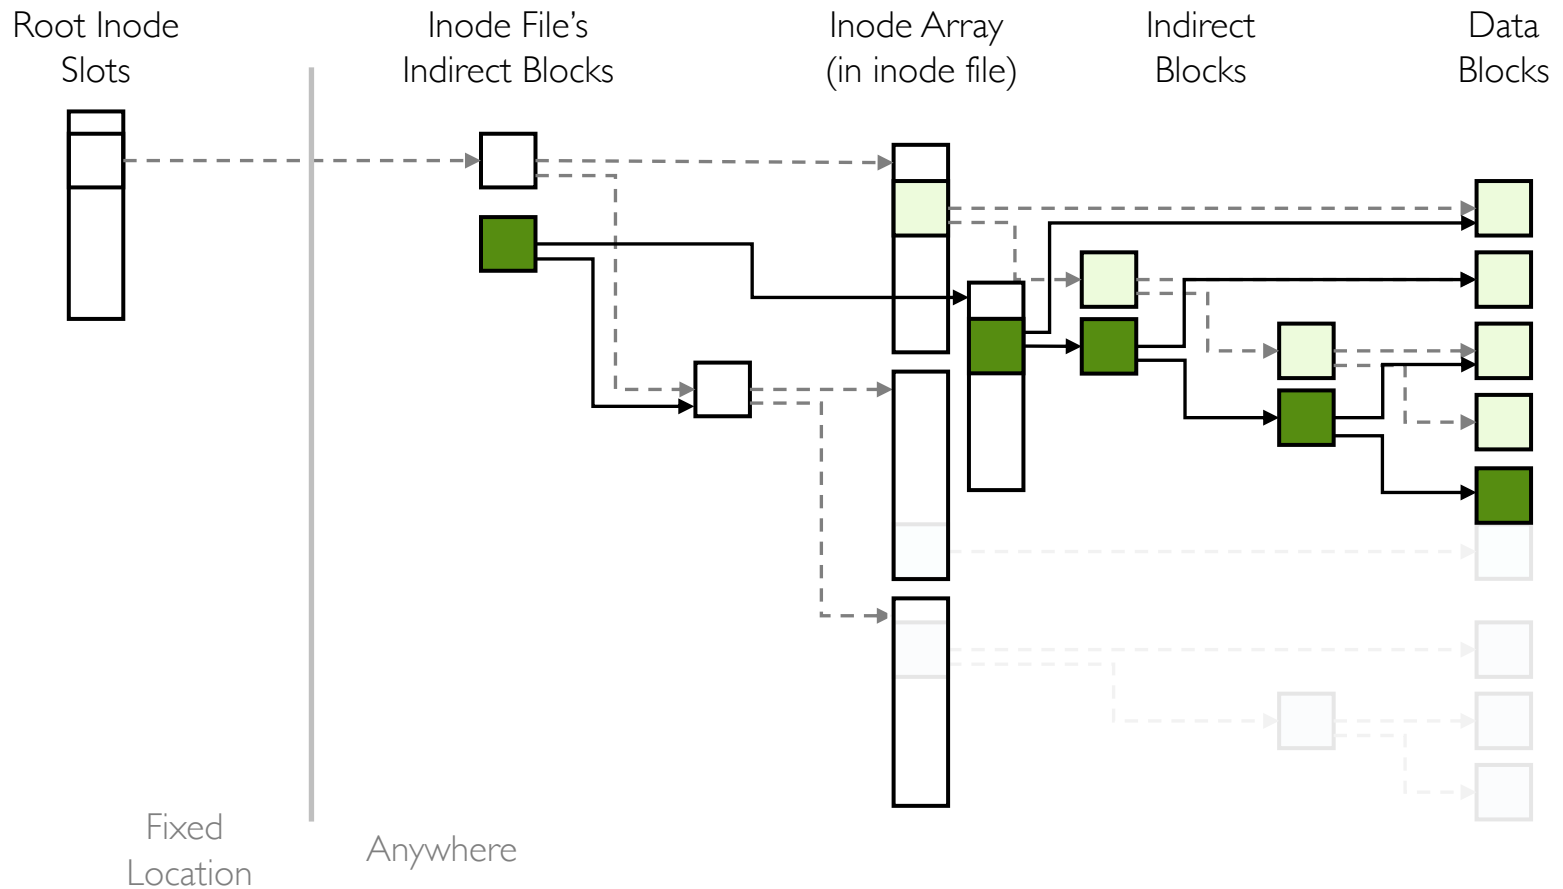
Simple Copy-on-write File System



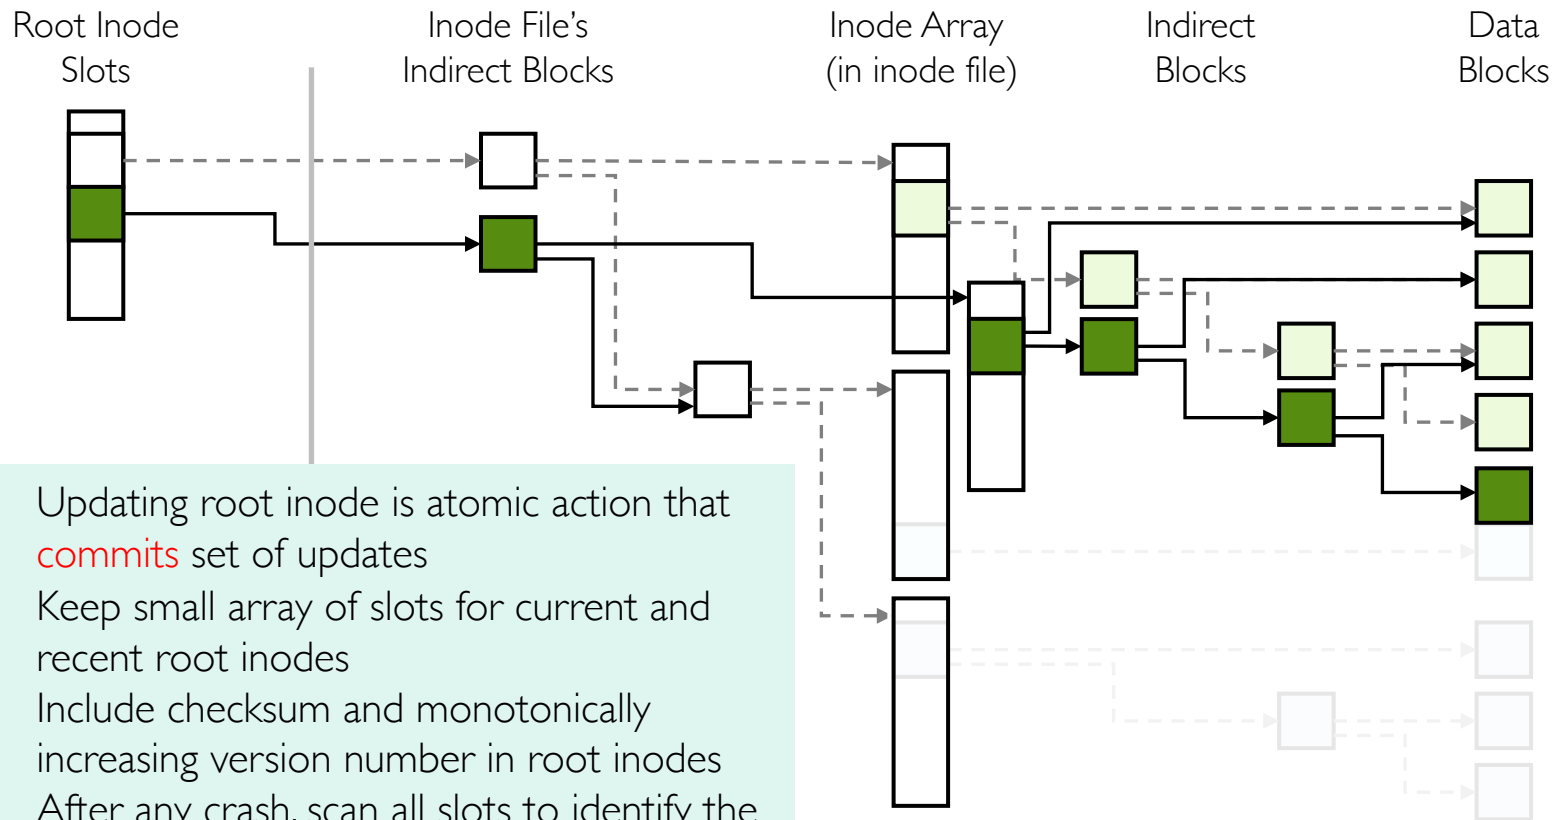
Simple Copy-on-write File System



Simple Copy-on-write File System

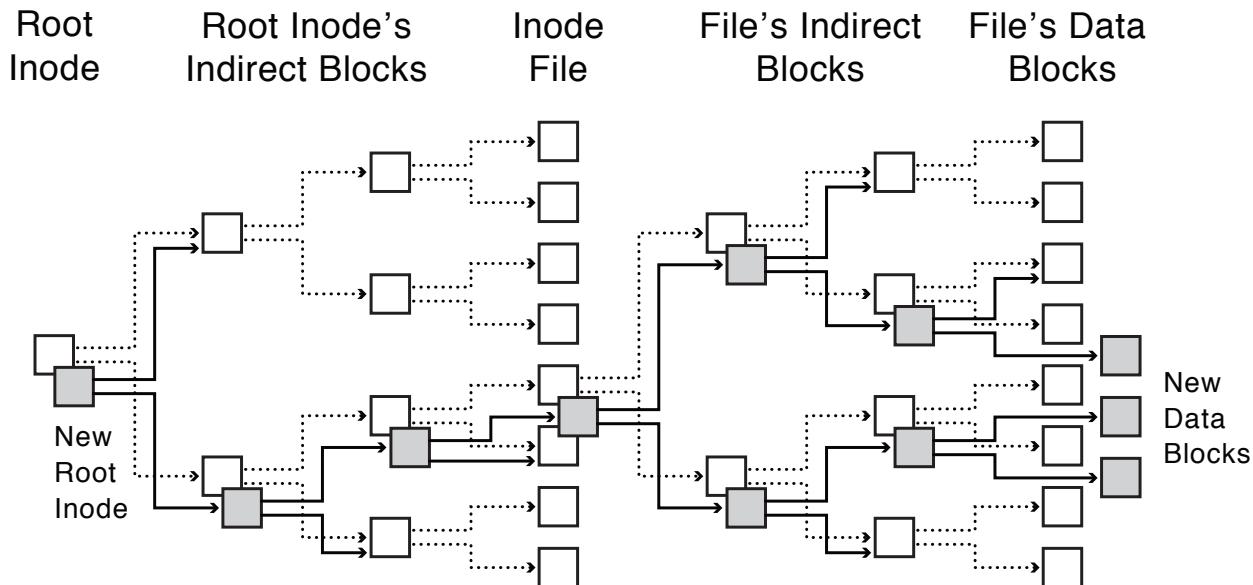


Simple Copy-on-write File System





- Updating root inode is atomic action that **commits** set of updates
- Keep small array of slots for current and recent root inodes
- Include checksum and monotonically increasing version number in root inodes
- After any crash, scan all slots to identify the newest root inode that includes correct checksum

Copy-on-write Batch Update



- Updates to file's inode and indirect blocks may be amortized over multiple writes to **different blocks of the file**
- Updates to the root inode and root inode's indirect blocks may be amortized over multiple writes to **different files**

Update-in-place vs. Copy-on-write File Systems



- Update-in-place file system updates data and metadata in their existing locations
- Copy-on-write file system makes **new copies** on every update
 - Updates can be grouped into single write to free range of sequential blocks on disk
 - Sequential writes are much faster than random ones → **excellent write performance** even though copy-on-write writes more data than update-in-place does

Copy-on-write File Systems: Garbage Collection

- Write efficiency requires contiguous sequences of free blocks
 - Updates leave dead blocks scattered
- Read efficiency requires data read together to be in the same block group
 - Write anywhere leaves related data scattered
- Solution: background coalescing of live/dead blocks

Copy-on-write File Systems: Discussion

- + Correct behavior regardless of failures
- + Fast recovery (root block array)
- + High throughput (best if updates are batched)
- – Potential for high latency
- – Small changes require many writes
- – Garbage collection essential for performance

Outline

- Problem posed by machine/disk failures
- Transaction concept (ASID properties)
- Transactional file systems
 - Journaling and logging
- Copy-on-write file systems
- Redundant arrays of independent disks (RAID)

How to Make File System Durable?


- Disk drives use error correcting codes (ECC) to deal with small defects
 - Can allow recovery of data from small media defects
- Writes should survive in **short term**
 - Either abandon delayed writes or
 - Use battery-backed RAM (non-volatile RAM - NVRAM) for dirty blocks in buffer cache
- Data should survive in **long term**
 - Need to replicate! Keep more than one copy of data!
 - Important element: independence of failure
 - Could put copies on one disk, but if disk head fails ...
 - Could put copies on different disks, but if server fails ...
 - Could put copies on different servers, but if building is struck by lightning ...
 - Could put copies on servers in different continents ...

RAID: Redundant Arrays of Inexpensive Disks

- Invented by David Patterson, Garth A. Gibson, and Randy Katz in 1987
- Data stored on multiple disks (redundancy)
- Implemented either in software or hardware
 - In hardware case, done by disk controller; file system may not even know that there is more than one disk in use
- Initially, five levels of RAID (more now!)


RAID 0: Striping

- Multiple drives (at least 2) are used to store data
- Data is split into blocks written across all drives
- + Offers great performance
 - Read and write operations can be parallelized over multiple drives
- + Uses entire capacity and is easy to implement
- – Not fault-tolerant
 - If one drive fails, data cannot be restored (not suitable for mission-critical systems)



RAID I: Mirroring

- Data is stored twice
- Each drive is fully duplicated onto its “shadow”
- + Offers low-overhead data recovery
 - Disk failure \Rightarrow replace disk & copy data to new disk
 - Hot spare: idle disk already attached to system to be used for immediate replacement
- – 100% capacity overhead
- – Bandwidth sacrificed on write
 - Logical write = two physical writes




RAID 5: Striping With Parity

- With G drives, for every $G - 1$ blocks, save one parity block (XOR of all data blocks)

- $P = D_1 \oplus D_2 \oplus D_3 \oplus \dots \oplus D_{G-1}$
- If D_1 is lost, it can be restored using P and other blocks: $D_1 = P \oplus D_2 \oplus \dots \oplus D_{G-1}$

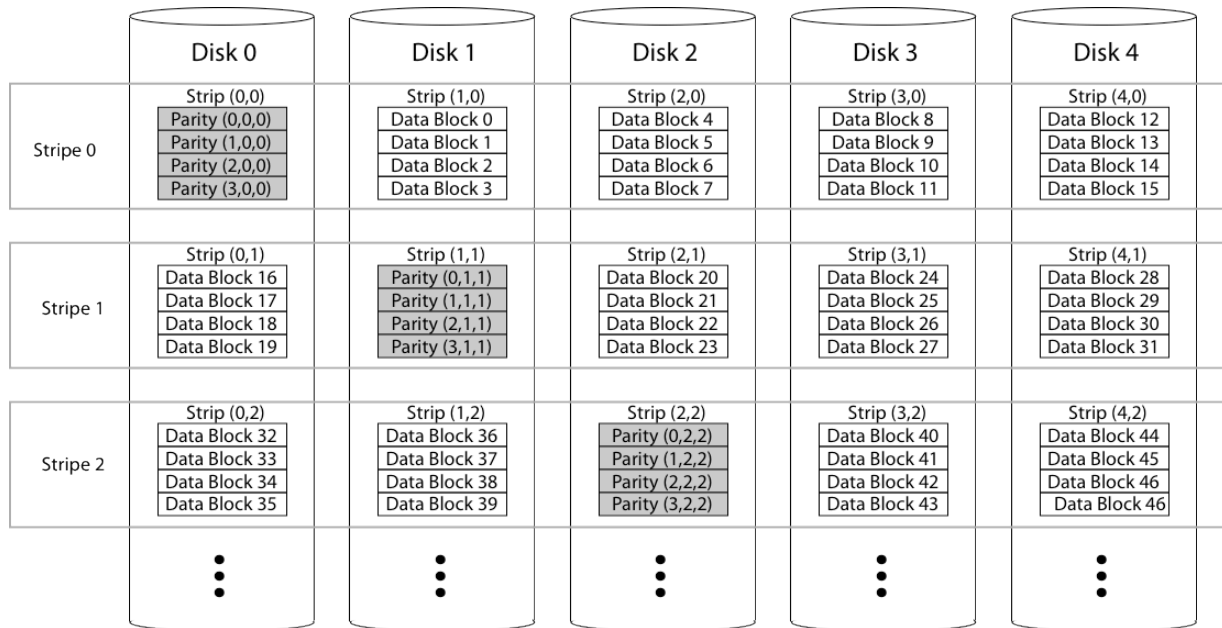
- Parity for given set of blocks must be updated each time any of blocks is updated

- To balance load, parity blocks are stored rotationally across all drives



- + Offers high availability with low capacity overhead
- – High recovery overhead
- – Complex implementation with high-overhead writes (requires 4 I/O operations)
 - Read old data
 - Read old parity
 - Write new data
 - Write new parity
 - Remove old data from parity ($P_{tmp} = P_{old} \oplus D_{old}$)
 - Calculate new parity ($P_{new} = P_{tmp} \oplus D_{new}$)


RAID 5: Closer Look



- To balance parallelism versus sequential-access efficiency, single **strip** of several sequential blocks is placed on one drive before shifting to another
- Set of $G - 1$ data strips and their parity strip is called **stripe**


RAID 6: Striping With Double Parity

- With RAID 5, time to repair failed drive is too long, another drive might fail in process!
- RAID 6 is like RAID 5, but with **two parity data**
 - Simple XOR will no longer work
 - We need something more complex
- Example: with G drives and k -bit blocks, if $G - 2 \leq k$
 - $P = D_1 \oplus D_2 \oplus D_3 \oplus \dots \oplus D_{G-2}$
 - $Q = D_1 \oplus \text{shift}(D_2) \oplus \text{shift}^2(D_3) \oplus \dots \oplus \text{shift}^{G-3}(D_{G-2})$
 - If D_1 is lost, recover it using P and other blocks (like RAID 5)
 - If D_1 and P are lost, recover it using Q and other blocks
 - $D_1 = Q \oplus \text{shift}(D_2) \oplus \text{shift}^2(D_3) \oplus \dots \oplus \text{shift}^{G-3}(D_{G-2})$
 - If D_1 and D_2 are lost, recover them using P, Q , and other blocks
 - $D_1 \oplus D_2 = P \oplus D_3 \oplus \dots \oplus D_{G-2}$
 - $D_1 \oplus \text{shift}(D_2) = Q \oplus \text{shift}^2(D_3) \oplus \dots \oplus \text{shift}^{G-3}(D_{G-2})$
 - System of $2k$ equations in $2k$ unknowns which uniquely determines the lost data
 - If $G - 2 > k$ or if we need more parity blocks, then general option is **Reed-Solomon system**
- + More reliable than RAID 5
 - Tolerates two drive failures
- – Even more slower writes than RAID 5
- – Higher capacity overhead than RAID 5



RAID 10: Combining RAID 1 & RAID 0


- RAID 10: RAID 1 + RAID 0
- + Advantages of RAID 1 and RAID 0
- – Higher capacity overhead than RAID 5 and 6



Summary

- Transactions
 - ACID properties: atomicity, consistency, isolation, and durability
 - Redo logging and two-phase locking
- Transactional file systems
 - Journaling and logging
- Copy-on-write file systems
 - Write new version of file system on each update
- Redundant arrays of inexpensive disks (RAID)
 - RAID 0, 1, 5, 6, and 10

Questions?



Acknowledgment

- Slides by courtesy of Anderson, Culler, Stoica, Silberschatz, Joseph, and Canny